

sta-ass-4

August 30, 2024

```
[ ]: #Question1
import pandas as pd
d=pd.read_csv('/train (1).csv')
d
y=d['price_range']
print(y)
```

```
0      1
1      2
2      2
3      2
4      1
..
1995    0
1996    2
1997    3
1998    0
1999    3
Name: price_range, Length: 2000, dtype: int64
```

```
[ ]: x=d.drop('price_range',axis=1)
print(x)
```

```
      battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  \
0           842      0         2.2          0    1      0           7
1          1021      1         0.5          1    0      1          53
2           563      1         0.5          1    2      1          41
3           615      1         2.5          0    0      0          10
4          1821      1         1.2          0   13      1          44
...           ...    ...           ...    ..    ...           ...
1995          794      1         0.5          1    0      1           2
1996         1965      1         2.6          1    0      0          39
1997         1911      0         0.9          1    1      1          36
1998         1512      0         0.9          0    4      1          46
1999          510      1         2.0          1    5      1          45

      m_dep  mobile_wt  n_cores  pc  px_height  px_width  ram  sc_h  sc_w  \
0         0.6       188        2   2         20       756  2549    9    7
```

1	0.7	136	3	6	905	1988	2631	17	3
2	0.9	145	5	6	1263	1716	2603	11	2
3	0.8	131	6	9	1216	1786	2769	16	8
4	0.6	141	2	14	1208	1212	1411	8	2
...
1995	0.8	106	6	14	1222	1890	668	13	4
1996	0.2	187	4	3	915	1965	2032	11	10
1997	0.7	108	8	3	868	1632	3057	9	1
1998	0.1	145	5	5	336	670	869	18	10
1999	0.9	168	6	16	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi
0	19	0	0	1
1	7	1	1	0
2	9	1	1	0
3	11	1	0	0
4	15	1	1	0
...
1995	19	1	1	0
1996	16	1	1	1
1997	5	1	1	0
1998	19	1	1	1
1999	2	1	1	1

[2000 rows x 20 columns]

```
[ ]: dd=(x-x.min())/(x.max()-x.min())
      print(dd)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	\
0	0.227789	0.0	0.68	0.0	0.052632	0.0	
1	0.347361	1.0	0.00	1.0	0.000000	1.0	
2	0.041416	1.0	0.00	1.0	0.105263	1.0	
3	0.076152	1.0	0.80	0.0	0.000000	0.0	
4	0.881764	1.0	0.28	0.0	0.684211	1.0	
...
1995	0.195725	1.0	0.00	1.0	0.000000	1.0	
1996	0.977956	1.0	0.84	1.0	0.000000	0.0	
1997	0.941884	0.0	0.16	1.0	0.052632	1.0	
1998	0.675351	0.0	0.16	0.0	0.210526	1.0	
1999	0.006012	1.0	0.60	1.0	0.263158	1.0	

	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width	\
0	0.080645	0.555556	0.900000	0.142857	0.10	0.010204	0.170895	
1	0.822581	0.666667	0.466667	0.285714	0.30	0.461735	0.993324	
2	0.629032	0.888889	0.541667	0.571429	0.30	0.644388	0.811749	
3	0.129032	0.777778	0.425000	0.714286	0.45	0.620408	0.858478	
4	0.677419	0.555556	0.508333	0.142857	0.70	0.616327	0.475300	

...
1995	0.000000	0.777778	0.216667	0.714286	0.70	0.623469	0.927904
1996	0.596774	0.111111	0.891667	0.428571	0.15	0.466837	0.977971
1997	0.548387	0.666667	0.233333	1.000000	0.15	0.442857	0.755674
1998	0.709677	0.000000	0.541667	0.571429	0.25	0.171429	0.113485
1999	0.693548	0.888889	0.733333	0.714286	0.80	0.246429	0.169559

	ram	sc_h	sc_w	talk_time	three_g	touch_screen	wifi
0	0.612774	0.285714	0.388889	0.944444	0.0	0.0	1.0
1	0.634687	0.857143	0.166667	0.277778	1.0	1.0	0.0
2	0.627205	0.428571	0.111111	0.388889	1.0	1.0	0.0
3	0.671566	0.785714	0.444444	0.500000	1.0	0.0	0.0
4	0.308658	0.214286	0.111111	0.722222	1.0	1.0	0.0

...
1995	0.110102	0.571429	0.222222	0.944444	1.0	1.0	0.0
1996	0.474613	0.428571	0.555556	0.777778	1.0	1.0	1.0
1997	0.748530	0.285714	0.055556	0.166667	1.0	1.0	0.0
1998	0.163816	0.928571	0.555556	0.944444	1.0	1.0	1.0
1999	0.978888	1.000000	0.222222	0.000000	1.0	1.0	1.0

[2000 rows x 20 columns]

```
[ ]: from sklearn.model_selection import train_test_split#for dividing test and train
```

```
[ ]: x_train,x_test,y_train,y_test=train_test_split(dd,y,test_size=0.
↪2,random_state=100)
```

```
[ ]: print(x_train)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	\
1260	0.398130	0.0	0.36	0.0	0.000000	0.0	
916	0.210421	1.0	0.00	0.0	0.210526	0.0	
532	0.184369	1.0	0.84	0.0	0.157895	0.0	
1159	0.002004	0.0	0.92	1.0	0.000000	0.0	
1584	0.983300	1.0	0.76	1.0	0.315789	0.0	
...
1879	0.321977	0.0	0.52	0.0	0.000000	1.0	
1895	0.566466	1.0	0.64	1.0	0.157895	0.0	
1859	0.098864	1.0	0.36	0.0	0.210526	1.0	
792	0.765531	0.0	0.36	1.0	0.157895	1.0	
1544	0.553774	1.0	0.52	1.0	0.052632	0.0	
	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width \
1260	0.935484	0.777778	0.841667	0.000000	0.85	0.206633	0.012684
916	0.193548	0.888889	0.983333	0.285714	0.65	0.076020	0.038718
532	0.451613	0.888889	0.566667	0.571429	0.85	0.096429	0.429239
1159	0.612903	0.444444	0.816667	0.285714	0.00	0.319388	0.463952
1584	0.709677	0.000000	0.941667	0.428571	0.50	0.356633	0.234312

...
1879	0.822581	0.888889	0.891667	0.000000	0.20	0.510714	0.451268
1895	0.709677	0.888889	0.450000	0.714286	0.60	0.108163	0.082109
1859	0.274194	0.777778	0.916667	0.857143	0.60	0.175510	0.701602
792	0.580645	0.222222	0.183333	0.714286	0.75	0.098469	0.457276
1544	0.612903	0.444444	0.808333	0.000000	0.25	0.098980	0.261682

	ram	sc_h	sc_w	talk_time	three_g	touch_screen	wifi
1260	0.067611	0.714286	0.666667	0.555556	1.0	0.0	0.0
916	0.203634	0.857143	0.555556	0.722222	1.0	0.0	0.0
532	0.970069	0.642857	0.555556	0.277778	1.0	0.0	0.0
1159	0.057189	0.071429	0.000000	0.777778	1.0	0.0	0.0
1584	0.102352	0.071429	0.000000	0.055556	0.0	0.0	1.0

...
1879	0.302779	0.857143	0.444444	0.333333	1.0	1.0	1.0
1895	0.564137	0.428571	0.388889	0.444444	1.0	0.0	1.0
1859	0.129075	0.928571	0.055556	0.000000	1.0	1.0	0.0
792	0.608231	0.785714	0.777778	0.722222	1.0	0.0	1.0
1544	0.647247	0.285714	0.222222	0.666667	1.0	1.0	1.0

[1600 rows x 20 columns]

```
[ ]: print(x_test)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	\
1025	0.387442	1.0	0.80	1.0	0.684211	0.0	
1208	0.325985	1.0	0.08	0.0	0.684211	1.0	
1055	0.814295	1.0	0.08	1.0	0.000000	0.0	
367	0.216433	0.0	0.96	1.0	0.000000	0.0	
815	0.675351	1.0	0.00	0.0	0.052632	0.0	
...
807	0.541750	0.0	0.32	0.0	0.263158	1.0	
711	0.830327	1.0	0.20	0.0	0.052632	1.0	
1541	0.215097	0.0	0.20	1.0	0.473684	1.0	
1001	0.683367	0.0	0.84	1.0	0.263158	0.0	
687	0.178357	0.0	0.00	1.0	0.000000	0.0	

	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width	\
1025	0.758065	1.000000	0.158333	0.285714	0.85	0.144898	0.012684	
1208	0.790323	0.333333	0.433333	0.285714	0.75	0.341837	0.202937	
1055	0.274194	0.000000	0.475000	0.142857	0.15	0.208673	0.097463	
367	0.516129	0.777778	0.158333	1.000000	0.85	0.750510	0.911883	
815	0.806452	0.000000	0.425000	0.000000	0.70	0.169898	0.301736	
...
807	0.129032	0.555556	0.691667	0.142857	0.45	0.048469	0.262350	
711	0.290323	0.444444	0.266667	1.000000	0.15	0.034694	0.145527	
1541	0.258065	0.111111	0.850000	0.571429	0.80	0.602041	0.567423	
1001	0.887097	0.444444	0.366667	0.142857	0.60	0.320918	0.090120	

```
687      0.274194  0.888889   0.658333  0.142857  0.05   0.325000  0.327770
```

	ram	sc_h	sc_w	talk_time	three_g	touch_screen	wifi
1025	0.018439	0.285714	0.055556	0.222222	0.0	0.0	1.0
1208	0.559327	0.357143	0.055556	1.000000	1.0	0.0	1.0
1055	0.258151	0.571429	0.111111	0.277778	0.0	0.0	0.0
367	0.710583	0.214286	0.111111	0.444444	1.0	1.0	0.0
815	0.733832	0.071429	0.277778	0.277778	0.0	0.0	1.0
...
807	0.984768	0.071429	0.111111	0.611111	1.0	1.0	1.0
711	0.253608	0.928571	0.388889	0.666667	1.0	0.0	0.0
1541	0.972207	0.571429	0.555556	0.388889	1.0	1.0	0.0
1001	0.184661	0.428571	0.222222	0.777778	0.0	1.0	1.0
687	0.237573	1.000000	0.055556	0.277778	0.0	0.0	1.0

```
[400 rows x 20 columns]
```

```
[ ]: print(y_train)
```

```
1260    0
916     0
532     3
1159    0
1584    0
..
1879    1
1895    1
1859    0
792     2
1544    2
Name: price_range, Length: 1600, dtype: int64
```

```
[ ]: print(y_test)
```

```
1025    0
1208    1
1055    1
367     3
815     2
..
807     3
711     1
1541    3
1001    0
687     0
Name: price_range, Length: 400, dtype: int64
```

```
[ ]: print(x.isnull())
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
...	
1995	False	False	False	False	False	False	False	
1996	False	False	False	False	False	False	False	
1997	False	False	False	False	False	False	False	
1998	False	False	False	False	False	False	False	
1999	False	False	False	False	False	False	False	

	m_dep	mobile_wt	n_cores	pc	px_height	px_width	ram	sc_h	\
0	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	
...	
1995	False	False	False	False	False	False	False	False	
1996	False	False	False	False	False	False	False	False	
1997	False	False	False	False	False	False	False	False	
1998	False	False	False	False	False	False	False	False	
1999	False	False	False	False	False	False	False	False	

	sc_w	talk_time	three_g	touch_screen	wifi
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...
1995	False	False	False	False	False
1996	False	False	False	False	False
1997	False	False	False	False	False
1998	False	False	False	False	False
1999	False	False	False	False	False

[2000 rows x 20 columns]

```
[ ]: print(d.dtypes)
```

```
battery_power    int64
blue              int64
clock_speed      float64
dual_sim         int64
fc               int64
```

```

four_g          int64
int_memory      int64
m_dep          float64
mobile_wt       int64
n_cores         int64
pc              int64
px_height       int64
px_width        int64
ram             int64
sc_h            int64
sc_w            int64
talk_time       int64
three_g         int64
touch_screen    int64
wifi            int64
price_range     int64
dtype: object

```

```

[8]: #2)Question
      # Read the data with pandas and describe the data
      import pandas as pd
      d=pd.read_csv("/content/sample_data/california_housing_train.csv")
      print(d)
      print(d.describe)

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-114.31	34.19	15.0	5612.0	1283.0	
1	-114.47	34.40	19.0	7650.0	1901.0	
2	-114.56	33.69	17.0	720.0	174.0	
3	-114.57	33.64	14.0	1501.0	337.0	
4	-114.57	33.57	20.0	1454.0	326.0	
...	
16995	-124.26	40.58	52.0	2217.0	394.0	
16996	-124.27	40.69	36.0	2349.0	528.0	
16997	-124.30	41.84	17.0	2677.0	531.0	
16998	-124.30	41.80	19.0	2672.0	552.0	
16999	-124.35	40.54	52.0	1820.0	300.0	

	population	households	median_income	median_house_value
0	1015.0	472.0	1.4936	66900.0
1	1129.0	463.0	1.8200	80100.0
2	333.0	117.0	1.6509	85700.0
3	515.0	226.0	3.1917	73400.0
4	624.0	262.0	1.9250	65500.0
...
16995	907.0	369.0	2.3571	111400.0
16996	1194.0	465.0	2.5179	79000.0
16997	1244.0	456.0	3.0313	103600.0

16998	1298.0	478.0	1.9797	85800.0
16999	806.0	270.0	3.0147	94600.0

[17000 rows x 9 columns]

```
<bound method NDFrame.describe of
total_rooms  total_bedrooms  \
0           -114.31      34.19      15.0      5612.0      1283.0
1           -114.47      34.40      19.0      7650.0      1901.0
2           -114.56      33.69      17.0       720.0       174.0
3           -114.57      33.64      14.0     1501.0       337.0
4           -114.57      33.57      20.0     1454.0       326.0
...          ...          ...          ...          ...          ...
16995       -124.26      40.58      52.0     2217.0       394.0
16996       -124.27      40.69      36.0     2349.0       528.0
16997       -124.30      41.84      17.0     2677.0       531.0
16998       -124.30      41.80      19.0     2672.0       552.0
16999       -124.35      40.54      52.0     1820.0       300.0
```

	population	households	median_income	median_house_value
0	1015.0	472.0	1.4936	66900.0
1	1129.0	463.0	1.8200	80100.0
2	333.0	117.0	1.6509	85700.0
3	515.0	226.0	3.1917	73400.0
4	624.0	262.0	1.9250	65500.0
...
16995	907.0	369.0	2.3571	111400.0
16996	1194.0	465.0	2.5179	79000.0
16997	1244.0	456.0	3.0313	103600.0
16998	1298.0	478.0	1.9797	85800.0
16999	806.0	270.0	3.0147	94600.0

[17000 rows x 9 columns]>

[9]: *#Find data type and shape of each column*

```
print(d.dtypes)
print(d.shape)
```

```
longitude      float64
latitude       float64
housing_median_age  float64
total_rooms    float64
total_bedrooms  float64
population     float64
households     float64
median_income  float64
median_house_value float64
dtype: object
(17000, 9)
```



```
[10]: #Find the target and features
x=d.drop('median_house_value',axis=1)
y=d['median_house_value']
print(x)
print(y)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-114.31	34.19	15.0	5612.0	1283.0	
1	-114.47	34.40	19.0	7650.0	1901.0	
2	-114.56	33.69	17.0	720.0	174.0	
3	-114.57	33.64	14.0	1501.0	337.0	
4	-114.57	33.57	20.0	1454.0	326.0	
...	
16995	-124.26	40.58	52.0	2217.0	394.0	
16996	-124.27	40.69	36.0	2349.0	528.0	
16997	-124.30	41.84	17.0	2677.0	531.0	
16998	-124.30	41.80	19.0	2672.0	552.0	
16999	-124.35	40.54	52.0	1820.0	300.0	

	population	households	median_income
0	1015.0	472.0	1.4936
1	1129.0	463.0	1.8200
2	333.0	117.0	1.6509
3	515.0	226.0	3.1917
4	624.0	262.0	1.9250
...
16995	907.0	369.0	2.3571
16996	1194.0	465.0	2.5179
16997	1244.0	456.0	3.0313
16998	1298.0	478.0	1.9797
16999	806.0	270.0	3.0147

[17000 rows x 8 columns]

0	66900.0
1	80100.0
2	85700.0
3	73400.0
4	65500.0
...	
16995	111400.0
16996	79000.0
16997	103600.0
16998	85800.0
16999	94600.0

Name: median_house_value, Length: 17000, dtype: float64

```
[11]: #Find the null values (if yes fill the null values with '0' or mean of that
      ↪ column)
      print(d.isnull())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	
...	
16995	False	False	False	False	False	
16996	False	False	False	False	False	
16997	False	False	False	False	False	
16998	False	False	False	False	False	
16999	False	False	False	False	False	

	population	households	median_income	median_house_value
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
16995	False	False	False	False
16996	False	False	False	False
16997	False	False	False	False
16998	False	False	False	False
16999	False	False	False	False

[17000 rows x 9 columns]

```
[12]: #Normalize all the features
      dd=(x-x.min())/(x.max()-x.min())
      print(dd)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	1.000000	0.175345	0.274510	0.147885	0.198945	
1	0.984064	0.197662	0.352941	0.201608	0.294848	
2	0.975100	0.122210	0.313725	0.018927	0.026847	
3	0.974104	0.116897	0.254902	0.039515	0.052142	
4	0.974104	0.109458	0.372549	0.038276	0.050435	
...	
16995	0.008964	0.854410	1.000000	0.058389	0.060987	
16996	0.007968	0.866100	0.686275	0.061869	0.081782	
16997	0.004980	0.988310	0.313725	0.070515	0.082247	
16998	0.004980	0.984060	0.352941	0.070384	0.085506	

16999	0.000000	0.850159	1.000000	0.047924	0.046400
-------	----------	----------	----------	----------	----------

	population	households	median_income
0	0.028364	0.077454	0.068530
1	0.031559	0.075974	0.091040
2	0.009249	0.019076	0.079378
3	0.014350	0.037000	0.185639
4	0.017405	0.042921	0.098281
...
16995	0.025337	0.060516	0.128081
16996	0.033381	0.076303	0.139170
16997	0.034782	0.074823	0.174577
16998	0.036296	0.078441	0.102054
16999	0.022506	0.044236	0.173432

[17000 rows x 8 columns]

```
[17]: #Split the data into train and test.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(dd,y,test_size=0.
↪2,random_state=100)
print(x_train)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
14730	0.215139	0.555792	0.862745	0.048768	0.060832	
9389	0.511952	0.408077	0.333333	0.056834	0.063935	
7662	0.595618	0.144527	0.607843	0.008699	0.015829	
9146	0.530876	0.297556	0.666667	0.043443	0.060366	
13485	0.238048	0.616366	0.254902	0.079953	0.105369	
...	
16304	0.184263	0.571732	0.607843	0.069514	0.064246	
79	0.875498	0.027630	0.274510	0.030816	0.050745	
12119	0.289841	0.647184	0.627451	0.028338	0.041899	
14147	0.227092	0.548353	0.745098	0.037380	0.042055	
5640	0.614542	0.153029	0.764706	0.044708	0.066729	

	population	households	median_income
14730	0.031643	0.057227	0.152825
9389	0.040864	0.064792	0.104943
7662	0.004877	0.016609	0.195928
9146	0.032008	0.056241	0.071433
13485	0.044676	0.108864	0.214045
...
16304	0.028308	0.066930	0.390567
79	0.028616	0.048841	0.061261
12119	0.021021	0.038645	0.058365
14147	0.017994	0.045387	0.250003
5640	0.035791	0.066436	0.146557

[13600 rows x 8 columns]

```
[18]: print(x_test)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
1559	0.710159	0.141339	0.294118	0.001423	0.000931	
5775	0.613546	0.143464	0.725490	0.105865	0.152855	
10247	0.436255	0.679065	0.450980	0.022512	0.028554	
5234	0.618526	0.175345	0.862745	0.094715	0.095903	
1416	0.714143	0.027630	0.568627	0.064690	0.060366	
...	
9213	0.526892	0.375133	0.372549	0.070674	0.085661	
15348	0.205179	0.568544	1.000000	0.083538	0.108163	
14256	0.225100	0.700319	0.529412	0.109846	0.138734	
5688	0.614542	0.131775	0.372549	0.048768	0.086127	
6548	0.604582	0.181722	0.549020	0.067405	0.121974	

	population	households	median_income
1559	0.001009	0.002138	0.146557
5775	0.103674	0.152278	0.167543
10247	0.010791	0.017596	0.184873
5234	0.047171	0.099655	0.291727
1416	0.028560	0.064463	0.231328
...
9213	0.041873	0.089952	0.082978
15348	0.041901	0.114290	0.184797
14256	0.072424	0.133037	0.126281
5688	0.047899	0.091268	0.065972
6548	0.052776	0.126295	0.145081

[3400 rows x 8 columns]

```
[19]: print(y_train)
```

14730	101400.0
9389	58000.0
7662	256300.0
9146	56500.0
13485	143500.0
...	
16304	349200.0
79	69400.0
12119	55900.0
14147	232500.0
5640	206300.0

Name: median_house_value, Length: 13600, dtype: float64

```
[20]: print(y_test)
```

```
1559    500001.0
5775    113600.0
10247   107200.0
5234    201000.0
1416    198500.0
```

...

```
9213     64600.0
15348    218200.0
14256     65500.0
5688    152500.0
6548    162100.0
```

```
Name: median_house_value, Length: 3400, dtype: float64
```